

Multi Agent Systems

- Lab 3 -

MDP Value and Policy Iteration Analysis

Slides adapted from Reinforcement Learning class by Vien Ngo,
University of Stuttgart, 2016

Markov decision process

- A reinforcement learning problem that satisfies the Markov property is called a Markov decision process, or MDP.
- $\text{MDP} = \{S, \mathcal{A}, \mathcal{T}, \mathcal{R}, \mathcal{P}_0, \gamma\}$.
 - S : consists of all possible states.
 - \mathcal{A} : consists of all possible actions.
 - \mathcal{T} : is a transition function which defines the probability $\mathcal{T}(s', s, a) = \text{Pr}(s'|s, a)$.
 - \mathcal{R} : is a reward function which defines the reward $\mathcal{R}(s, a)$.
 - \mathcal{P}_0 : is the probability distribution over initial states.
 - $\gamma \in [0, 1]$: is a discount factor.

State Value Function

The **value** (*expected discounted return* – for infinite horizon settings) of a policy π when started in state s :

$$V^\pi(s) = \mathbb{E}_\pi \{ r_0 + \gamma r_1 + \gamma^2 r_2 + \cdots \mid s_0 = s \}$$

where $0 < \gamma < 1$ is the *discounting* factor

Optimality: policy π^* is optimal iff

$$\forall_s : V^{\pi^*}(s) = V^*(s) \quad \text{where} \quad V^*(s) = \max_{\pi} V^\pi(s)$$

Bellman Principle of Optimality

$$V^*(s) = \max_a \left[R(a, s) + \gamma \sum_{s'} P(s' | a, s) V^*(s') \right]$$
$$\pi^*(s) = \operatorname{argmax}_a \left[R(a, s) + \gamma \sum_{s'} P(s' | a, s) V^*(s') \right]$$

Value Iteration

Given the Bellman equation

$$V^*(s) = \max_a \left[R(a, s) + \gamma \sum_{s'} P(s' | a, s) V^*(s') \right]$$

→ iterate

$$\forall_s : V_{k+1}(s) = \max_a \left[R(a, s) + \gamma \sum_{s'} P(s' | \pi(s), s) V_k(s') \right]$$

stopping criterion:

$$\max_s |V_{k+1}(s) - V_k(s)| \leq \epsilon$$

Policy Iteration

Value Iteration computes V^* directly

To **evaluate** a given policy π , one needs to compute V^π , that is iterate using π instead of \max_a

$$\forall_s : V_{k+1}(s) = R(\pi(s), s) + \gamma \sum_{s'} P(s'|\pi(s), s) V_k(s')$$

Optimal policy can then be computed in iterative way

Policy Iteration

- 1) Initialise π_0 in a given way (e.g. randomly)
- 2) Iterate
 - Policy Evaluation: compute V^{π_k}
 - Policy Improvement: $\pi_{k+1}(s) = \underset{a \in A}{\operatorname{argmax}} [R(a, s) + \gamma \sum_{s'} P(s'|a, s) V^{\pi_k}(s)]$

Value Iteration variants

Gauss-Seidel Value Iteration

- Standard VI algorithm updates all states at next iteration using **old** values at previous iteration (each iteration finishes when all states get updated).

Algorithm 1 Standard Value Iteration Algorithm

```
1: while (!converged) do  
2:    $V_{old} = V$   
3:   for (each  $s \in \mathcal{S}$ ) do  
4:      $V(s) = \max_a \{R(s, a) + \gamma \sum_{s'} P(s'|s, a) V_{old}(s')\}$ 
```

- Gauss-Seidel VI updates each state using values from previous computation.

Algorithm 2 Gauss-Seidel Value Iteration Algorithm

```
1: while (!converged) do  
2:   for (each  $s \in \mathcal{S}$ ) do  
3:      $V(s) = \max_a \{R(s, a) + \gamma \sum_{s'} P(s'|s, a) V(s')\}$ 
```

Prioritised Sweeping

- Similar to Gauss-Seidel VI, but the sequence of states in each iteration is proportional to their update magnitudes (Bellman errors).
- Define Bellman error as $E(s; V_t) = |V_{t+1}(s) - V_t(s)|$ that is the change of s 's value after the most recent update.

Algorithm 3 Prioritised Sweeping VI Algorithm

- 1: Initialize $V_0(s)$ and priority values $H_0(s)$, $\forall s \in \mathcal{S}$.
 - 2: **for** $k = 0, 1, 2, 3, \dots$ **do**
 - 3: pick a state to update (with the highest priority): $s_k \in \arg \max_{s \in \mathcal{S}} H_k(s)$
 - 4: value update: $V_{k+1}(s_k) = \max_{a \in \mathcal{A}} [R(s_k, a_k) + \gamma \sum_{s'} P(s'|s_k, a_k) V_k(s')]$
 - 5: for $s \neq s_k$: $V_{k+1}(s) = V_k(s)$
 - 6: update priority values: $\forall s \in \mathcal{S}, H_{k+1}(s) \leftarrow E(s; V_{k+1})$ (Note: the error is w.r.t the future update).
-

OpenAI Gym Test Environments

- Three test MDP environments based on **OpenAI Gym**:
 - **Taxi-v2**
 - 4 locations
 - Pickup passenger at one location and drop him off at another
 - 6 actions: move NORTH, SOUTH, EAST, WEST + PICK_UP + DROP_OFF
 - Rewards: +20 for successful drop-off, -1 per movement, -10 for illegal pick-up or drop-off
 - **FrozenLake-v0 (small) and FrozenLake8x8-v0 (larger)**
 - Agent controls movement over a grid
 - Some tiles walkable, some lead agent to fall into water; agent has a start and goal tile
 - Movement of agent with uncertainty (due to slippery ice)
 - Rewards: +1 if agent finds correct path, 0 otherwise

Tasks

- For each game consider the following values for convergence criteria
 - max_iterations: $5 \cdot 10^5$
 - epsilon_threshold: 10^{-2} or 10^{-3}
 - Discount factor: 0.9
- Run the three variants of Value Iteration (**VI**, **Gauss-Seidel VI**, **Prioritized Sweeping VI**)
- Run 10 instantiations (random policy init each time) for **Policy Iteration** for each game
 - Compute averages of **number of iterations** until convergence
- **Plot convergence graph**
 - X axis: number of iterations (**NOTE!** An **iteration** is considered an **update to a state** in the value function, **i.e. an update to $V(s)$**)
 - Y axis: $\|V - V^*\|_2$
- Analyze convergence speed properties